

The Problem: A new treatment is given to 100 patients. Of them, only 8 respond. But there is a subgroup of 5 in which 3 patients respond, yielding a response rate of 60%! Should the treatment be recommended (or at least developed further) for people in the subgroup?

	Dark Hair (D)	Light Hair (L)	TOTAL
Responder (R)	3	5	8
Nonresponder (N)	2	90	92
TOTAL	5	95	100

We want to express prior knowledge or belief concerning how related $P_{R,D}$ and $P_{R,L}$ are.

Dr. Lump is sure that they are identical; Dr. Lump lumps together the D and L groups.

Dr. Split is sure that they are totally unrelated.

Others may have intermediate positions.

Or, how close to splitting or lumping seem appropriate will depend on the nature of X .

Let's transform the two conditional probabilities with the log odds function.

For $X = D$ or L ,

$$z_{R,X} = \text{logit}(\Pr(R | X)) = \log\left(\frac{\Pr(R | X)}{1 - \Pr(R | X)}\right)$$

Then a broad family of useful priors is bivariate normal in the $(\text{logit}(P_{R,D}), \text{logit}(P_{R,L}))$ plane:

$$(z_{R,D}, z_{R,L})^T \sim N(\mu_z, \Sigma_z) \text{ where } \mu_z = (\mu_{R,D}, \mu_{R,L})^T \text{ and } \Sigma_z = \begin{pmatrix} \tau + \phi & \tau \\ \tau & \tau + \phi \end{pmatrix}.$$

The prior means for the logits are $\mu_{R,D}, \mu_{R,L}$. The superscript " T " means "transpose"; customarily these vectors are written as column vectors. (We usually set the prior means of $z_{R,D}$ and $z_{R,L}$ to be the same, so $\mu_{R,D} = \mu_{R,L}$.) Any factors that might affect both response probabilities are represented by the covariance τ . Any that are not shared are represented by $\phi = (\tau + \phi) - \tau$. The prior correlation between $z_{R,D}$ and $z_{R,L}$ is $\tau / (\tau + \phi)$.

So the joint density of $(\text{logit}(P_{R,D}), \text{logit}(P_{R,L}))$ is

$$f_z(z_{R,D}, z_{R,L}) = (2\pi)^{-1/2} \det(\Sigma_z)^{-1/2} \exp\left(-\frac{1}{2}(z - \mu_z)^T \Sigma_z^{-1}(z - \mu_z)\right)$$

which can be converted to a density on the original scale by multiplying by the Jacobean of the logit transformation:

$$\begin{aligned}
f_P(P_{R,D}, P_{R,L}) &= f_z(z_{R,D}, z_{R,L}) \times \det \begin{pmatrix} \partial z_{R,D} / \partial P_{R,D} & 0 \\ 0 & \partial z_{R,L} / \partial P_{R,L} \end{pmatrix}^{-1} \\
&= f_z(\text{logit } P_{R,D}, \text{logit } P_{R,L}) \times P_{R,D}(1 - P_{R,D}) \times P_{R,L}(1 - P_{R,L})
\end{aligned}$$

because the Jacobean of the transformation is

$$\begin{aligned}
\frac{d}{dP} z &= \frac{d}{dP} (\log((P / (1 - P)))) = \frac{1}{P / (1 - P)} \frac{d}{dP} (1 - 1 / (1 - P)) \\
&= \frac{1 - P}{P} (1 - P)^{-2} = P^{-1} (1 - P)^{-1} = P^{-1} + (1 - P)^{-1}
\end{aligned}$$

We use $f_P(P_{R,D}, P_{R,L})$ to draw the contours of prior and posterior distributions for \mathbf{P} .

The model (conditional distribution of the data):

The observed proportions are also converted to logits: $x_{R,X} = \text{logit } \hat{P}_{R,X}$, and the variances of these logits is estimated with the delta method applied to the logit function.

$$\begin{aligned}
\text{var}(\text{logit}(\hat{\mathbf{P}})) &\doteq \text{var}(\hat{\mathbf{P}}) \left(\frac{d}{d\mathbf{P}} \text{logit } \mathbf{P} \right)^2 \\
&= \mathbf{n}^{-1} \mathbf{P} (1 - \mathbf{P}) \left(\mathbf{P}^{-1} (1 - \mathbf{P})^{-1} \right)^2 = \mathbf{n}^{-1} \left(\mathbf{P}^{-1} (1 - \mathbf{P})^{-1} \right) \\
&\doteq \mathbf{n}^{-1} \left(\hat{\mathbf{P}}^{-1} (1 - \hat{\mathbf{P}})^{-1} \right)
\end{aligned}$$

When there is a zero count, the variance formula yields infinity.

A small "fudge factor" can be added to each count, to produce a finite variance.

We model the data, conditional on the unknown $z_{R,D}, z_{R,L}$ (which is what we are interested in knowing), with a normal approximation for the logits:

$$\begin{pmatrix} x_{R,D} \\ x_{R,L} \end{pmatrix} \Bigg| \begin{pmatrix} z_{R,D} \\ z_{R,L} \end{pmatrix} \sim N \left(\begin{pmatrix} z_{R,D} \\ z_{R,L} \end{pmatrix}, \Sigma_{x|z} \right)$$

where $\Sigma_{x|z} = \begin{pmatrix} n^{-1} P_{R,D}^{-1} (1 - P_{R,D})^{-1} & 0 \\ 0 & n^{-1} P_{R,L}^{-1} (1 - P_{R,L})^{-1} \end{pmatrix}$.

(Notice the very tall "|" indicating conditioning on the z 's.) Conditionally, the data are independent, as we see from the off-diagonal zeros.

The joint distribution

The joint distribution of the unknown z 's and the observed x 's is the product of the prior distribution and the model. This is a good spot to simplify our notation.

$$Y_1 = \begin{pmatrix} z_{R,D} \\ z_{R,L} \end{pmatrix}, \quad Y_2 = \begin{pmatrix} x_{R,D} \\ x_{R,L} \end{pmatrix}, \quad \mu_1 = E(Y_1) = \mu_2 = E(Y_2) = \begin{pmatrix} \mu_{R,D} \\ \mu_{R,L} \end{pmatrix}$$

$$\Sigma_{11} = \text{var}(Y_1) = \Sigma_z = \begin{pmatrix} \tau + \phi & \tau \\ \tau & \tau + \phi \end{pmatrix}, \quad \Sigma_{22} = \text{var}(Y_2) = \Sigma_x = \Sigma_z + \Sigma_{x|z}$$

$$\Sigma_{12} = \Sigma_{21} = \text{cov}(Y_1, Y_2) = \Sigma_z$$

Then the joint distribution of the Y 's (in our case, z and x) is

$$\begin{pmatrix} Y_1 \\ Y_2 \end{pmatrix} \sim N \left(\begin{pmatrix} \mu_1 \\ \mu_2 \end{pmatrix}, \begin{pmatrix} \Sigma_{11} & \Sigma_{12} \\ \Sigma_{21} & \Sigma_{22} \end{pmatrix} \right)$$

and the beautiful thing is that the posterior distribution we want is the conditional distribution of the unknown Y_1 (which is z) conditional on the observed Y_2 which is (x) . We get this from the standard formulas for conditional mean and variance for the multivariate normal distribution:

$$Y_2 | Y_1 = y_1 \sim N(\mu_{2.1}, \Sigma_{22.1})$$

where

$$\mu_{2.1} = E(Y_2 | Y_1 = y_1) = \mu_2 + \Sigma_{21} \Sigma_{11}^{-1} (y_1 - \mu_1)$$

and

$$\Sigma_{22.1} = \text{Var}(Y_2 | Y_1 = y_1) = \Sigma_{22} - \Sigma_{21} \Sigma_{11}^{-1} \Sigma_{12}.$$

Here it is in code. It's not hard to do. The function `solve()` does the matrix inversion.

```
varhat = apply(DLdata, 1, function(r) sum(1/r))
sig11 = sig12 = sig21 = matrix(c(tau+phi, tau, tau, tau+phi), nrow=2)
sig22 = sig11 + diag(varhat) ## marginal variance of the data
logit.prior.mean = c(mu0, mu0)
## always use fudged here:
postmean.logit = logit.prior.mean +
  sig12%%solve(sig22) %% (logit.hat.fudged-logit.prior.mean)
postmean.p = antilogit(postmean.logit)
postvar.logit = sig11 - sig12%%solve(sig22)%%sig21
```

To draw the contour plots, here is the code:

<https://github.com/professorbeautiful/Tl5lumpsplit/blob/master/inst/Tl5lumpsplit/Plight-Pdark-posterior-new.R>